

B-Tree

1

B-Tree

- Original B-Tree proposed by R. Bayer and E. McCreigh in 1972.
- A B-Tree is a specialized multi-way tree designed especially for use on external disk.
- Improved versions of B-Trees were later proposed in 1982 by Huddleston and Mehlhorn, and by Maier and Salveter.
- B-tree variants are used mostly today as index structures in database applications.

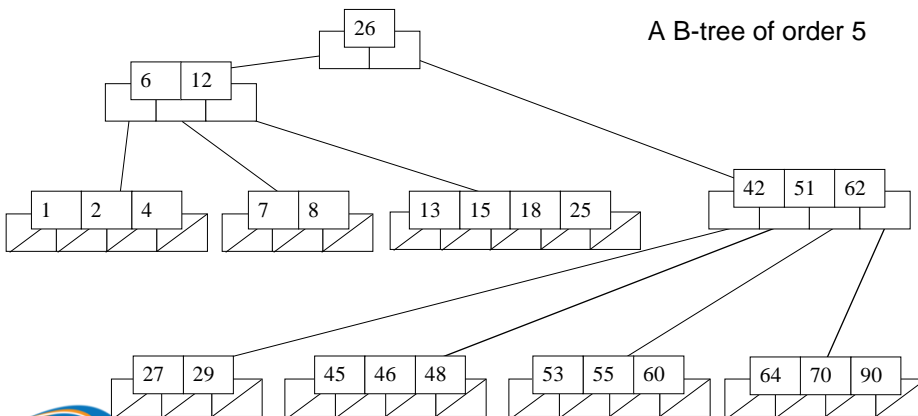
2

Definition

- (Knuth's definition) A B-tree of order m is a tree which satisfies the following properties:
 - Every node has at most m children.
 - Every non-leaf node (except root) has at least $\lceil m/2 \rceil$ child nodes.
 - The root has **at least two children** if it is not a leaf node.
 - A non-leaf node with k children contains $k - 1$ keys.
 - All leaves appear in the same level and carry no information.
- The number m should be (always) odd.


Example

A B-tree of order 5



Example

Is this a B-tree?



Height of B-Tree

- The maximum number of keys in a B-tree of order m and height h :

root	$m - 1$
level 2	$m(m - 1)$
level 3	$m^2(m - 1)$
...	
level h	$m^{h-1}(m - 1)$

- So, the total number of keys is

$$(1 + m + m^2 + m^3 + \dots + m^{h-1})(m - 1) =$$

$$[(m^h - 1) / (m - 1)] (m - 1) = \mathbf{m^h - 1}$$

Operations

- Insert (a key)
- Remove (a key)

Insertion

Insertion

- B-tree insertion is a generalization of 2-3 tree insertion.
- Insert K into B-tree of order m .
 - We find the insertion point (in a leaf) by doing a search.
 - If there is room then insert K .
 - Else, promote the middle key to the parent, split the node into nodes around the middle key.
- If the splitting backs up to the root, then
 - Make a new root containing the middle key.
- Note
 - The tree grows from the leaves, balance is always maintained.

Example

- Perform step-by-step the insertion of the following values to the initially empty B-tree order 5

1, 12, 8, 2, 25, 5, 14, 28, 17



Example



Insert 1

Insert 12



Insert 8

Insert 2

13



Example




Insert 25

14



Example




Insert 5, 14, 28

15



Example



Insert 17

16

Deletion

Deletion


- If the key to be deleted is not in a leaf, swap it with its successor (or predecessor). Then delete the key from the leaf.
- If leaf contains more than the minimum number of keys, then one can be deleted with no further action.

Deletion

- If the node contains the minimum number of keys, consider the two immediate siblings of the parent node:
 - If one of these siblings has more than the minimum number of keys, then **redistribute** one key from this sibling to the parent node, and one key from the parent to the deficient node.
 - If both immediate siblings have exactly the minimum number of keys, then **merge** the deficient node with one of the immediate sibling node and one key from the parent node.
- If this leaves the parent node with too few keys, then the process is propagated upward.


Example

Given a B-tree order 5:



Example


Given a B-tree order 5:



After removing 2

Example


Given a B-tree order 5:



Remove 52

Example


Given a B-tree order 5:



Remove 52 (replaced by 56)


Example

Given a B-tree order 5:



Example


Given a B-tree order 5:



Remove 72

Example


Given a B-tree order 5:



Not enough key. Merge nodes.


Example

Given a B-tree order 5:




Example

Given a B-tree order 5:



Example


Given a B-tree order 5:



(Redistribute) Borrow key from neighboring nodes


Example

Given a B-tree order 5:



Example

Given a B-tree order 5:




Remove 8



Merge 2 child nodes of 8

Example


Given a B-tree order 5:



Not enough keys at node 3.

Example

Given a B-tree order 5




A new root node

Decrease height of the tree

Quiz

Given the following B-tree order 5

- Remove 28, then remove 48



B-Trees & Efficiency

- Used in Mac, NTFS, OS2 for file structure.
- Allow insertion and deletion into a tree structure, based on $\log_m n$ property, where m is the order of the tree.
- The idea is that you leave some key spaces open. An insertion of a new key is done using available space (most cases).
 - Less dynamic than our typical Binary Tree
 - Ideal for disk-based operations.

B-Trees & Efficiency

- In practical applications, B-Trees of large order (e.g., $m = 128$) are more common than low-order B-Trees such as 2-3 trees.

Questions and Answers

fit@hcmus | DSA | 2024

38